

Utility nmon

Студент гр. 13541/3: Медведев М.А.

Преподаватель: Душутина Е.В.

nmon - системный администратор, тюнер, инструмент сравнения



For help type H or ...
nmon -? - hint
nmon -h - full details

To stop nmon type q to Quit

Fedora release 26 (Twenty Six) VERSION='26 (Workstation Edition)'
Vendor=GenuineIntel Model=Intel(R) Core(TM) i7-3520M CPU @ 2.90GHz
MHz=3392.328 bogomips=5786.79 lscpu:CPU=0 ttle Endian
ProcessorChips=1 PhysicalCores=2 Sockets=0 Cores=0 Thrds=0
Hyperthreads =2 VirtualCPUs =4 MHz=65 max=0 min=0

Use these keys to toggle statistics on/off:

c = CPU	l = CPU Long-term	- = Faster screen updates
C = " WideView	U = Utilisation	+ = Slower screen updates
m = Memory	V = Virtual memory	j = File Systems
d = Disks	n = Network	. = only busy disks/procs
r = Resource	N = NFS	h = more options
k = Kernel	t = Top-processes	q = Quit

Ключи

HELP: Hit h to remove this Info Hit q to Quit

Letters which toggle on/off statistics:

h = This help	r = Resources OS & Proc
c = CPU Util C = wide view	l = longer term CPU averages
m = Memory & Swap L=Huge	V = Virtual Memory
n = Network	N = NFS
d = Disk I/O Graphs D=Stats	o = Disks %Busy Map
k = Kernel stats & loadavg	j = Filesystem Usage J=reduced
M = MHz by thread & CPU	
t = TopProcess l=Priority/Nice/State	u = TopProc with command line
ReOrder by: 3=CPU 4=RAM 5=I/O	Hit u twice to update
g = User Defined Disk Groups	G = with -g switches Disk graphs
[start nmon with -g <filename>]	to disk groups only
	b = black & white mode
Other Controls:	
+ = double the screen refresh time	0 = reset peak marks (">") to zero
- = half the screen refresh time	space refresh screen now
. = Display only busy disks & CPU	q = Quit

(C) Copyright 2009 Nigel Griffiths

See <http://nmon.sourceforge.net>

Colour: #0# #1# #2# #3# #4# #5# #6# #7# #8# #9# #10# #11# #12#

Отображение информации с ключами -m & -r

Memory and Swap					
PageSize:4KB	RAM-Memory	Swap-Space	High-Memory	Low-Memory	
Total (MB)	7836.0	4096.0	- not in use	- not in use	
Free (MB)	1621.1	4096.0			
Free Percent	20.7%	100.0%			
Linux Kernel Internal Memory (MB)					
	Cached=	2350.8	Active=	3925.3	
Buffers=	557.7	Swapcached=	0.0	Inactive =	1611.3
Dirty =	2.8	Writeback =	0.0	Mapped =	569.7
Slab =	451.6	Commit_AS =	11393.4	PageTables=	80.2

```
Resources Linux & Processor
Linux: Linux version 4.15.4-200.fc26.x86_64 (mockbuild@bkernel02.phx2.fedoraproject.org)
Build: (gcc version 7.3.1 20180130 (Red Hat 7.3.1-2) (GCC))
Release : 4.15.4-200.fc26.x86_64
Version : #1 SMP Mon Feb 19 19:43:32 UTC 2018
cpuinfo: Vendor=GenuineIntel Model=Intel(R) Core(TM) i7-3520M CPU @ 2.90GHz
cpuinfo: Hz=2858.907 bogomips=5786.79
cpuinfo: ProcessorChips=1 PhyscalCores=2
cpuinfo: Hyperthreads =2 VirtualCPUs =4
# of CPUs: 4
Machine : x86_64
Nodename : localhost.localdomain
/etc/*ease[1]: Fedora release 26 (Twenty Six)
/etc/*ease[2]: NAME=Fedora
/etc/*ease[3]: VERSION="26 (Workstation Edition)"
/etc/*ease[4]: ID=fedora
lsb_release: Distributor ID: Fedora
lsb_release: Description: Fedora release 26 (Twenty Six)
lsb_release: Release: 26
lsb_release: Codename: TwentySix
```

Структура proc

```
534  int reread = 0;
535  struct {
536      FILE *fp;
537      char *filename;
538      int size;
539      int lines;
540      char *line[PROC_MAXLINES];
541      char *buf;
542      int read_this_interval;
543  } proc[P_NUMBER];
```


Откуда получаем данные ?

```
221 #define P_CPUINFO 0
222 #define P_STAT 1
223 #define P_VERSION 2
224 #define P_MEMINFO 3
225 #define P_UPTIME 4
226 #define P_LOADAVG 5
227 #define P_NFS 6
228 #define P_NFSD 7
229 #define P_VMSTAT 8 /* new in 13h */
230 #define P_NUMBER 9 /* one more than the max */
```

```
549 void proc_init()
550 {
551     proc[P_CPUINFO].filename = "/proc/cpuinfo";
552     proc[P_STAT].filename = "/proc/stat";
553     proc[P_VERSION].filename = "/proc/version";
554     proc[P_MEMINFO].filename = "/proc/meminfo";
555     proc[P_UPTIME].filename = "/proc/uptime";
556     proc[P_LOADAVG].filename = "/proc/loadavg";
557     proc[P_NFS].filename = "/proc/net/rpc/nfs";
558     proc[P_NFSD].filename = "/proc/net/rpc/nfsd";
559     proc[P_VMSTAT].filename = "/proc/vmstat";
560 }
```

Некоторые основные структуры данных для хранения статистики

```
1050 struct cpu_stat {
1051     long long user;
1052     long long nice;
1053     long long sys;
1054     long long idle;
1055     long long wait;
1056     long long irq;
1057     long long softirq;
1058     long long steal;
1059     long long guest;
1060     long long guest_nice;
1061     long long intr;
1062     long long ctxt;
1063     long long btime;
1064     long long procs;
1065     long long running;
1066     long long blocked;
1067     float uptime;
1068     float idletime;
1069     float mins1;
1070     float mins5;
1071     float mins15;
1072 };
```

```
1075 #define ulong unsigned long
1076 struct dsk_stat {
1077     char dk_name[32];
1078     int dk_major;
1079     int dk_minor;
1080     long dk_noinfo;
1081     ulong dk_reads;
1082     ulong dk_rmerge;
1083     ulong dk_rmsec;
1084     ulong dk_rkb;
1085     ulong dk_writes;
1086     ulong dk_wmerge;
1087     ulong dk_wmsec;
1088     ulong dk_wkb;
1089     ulong dk_xfers;
1090     ulong dk_bsize;
1091     ulong dk_time;
1092     ulong dk_inflight;
1093     ulong dk_backlog;
1094     ulong dk_partition;
1095     ulong dk_blocks;
1096     ulong dk_use;
1097     ulong dk_aveq;
1098 };
```

Некоторые основные структуры данных для хранения статистики

```
1100 struct mem_stat {
1101     long memtotal;
1102     long memfree;
1103     long memshared;
1104     long buffers;
1105     long cached;
1106     long swpcached;
1107     long active;
1108     long inactive;
1109     long hightotal;
1110     long highfree;
1111     long lowtotal;
1112     long lowfree;
1113     long swaptotal;
1114     long swapfree;
1115     #ifdef LARGEMEM
1116         long dirty;
1117         long writeback;
1118         long mapped;
1119         long slab;
1120         long committed_as;
1121         long pagetables;
1122         long hugetotal;
1123         long hugefree;
1124         long hugesize;
1125     #else
1126         long bigfree;
1127     #endif /*LARGEMEM*/
1128 };
```

```
1130 struct vm_stat {
1131     long long nr_dirty;
1132     long long nr_writeback;
1133     long long nr_unstable;
1134     long long nr_page_table_pages;
1135     long long nr_mapped;
1136     long long nr_slab;
1137     long long pgpgin;
1138     long long pgpgout;
1139     long long pswpin;
1140     long long pswpout;
1141     long long pgalloc_high;
1142     long long pgalloc_normal;
1143     long long pgalloc_dma;
1144     long long pgfree;
1145     long long pgactivate;
1146     long long pgdeactivate;
1147     long long pgfault;
1148     long long pgmajfault;
1149     long long pgrefill_high;
1150     long long pgrefill_normal;
1151     long long pgrefill_dma;
1152     long long pgsteal_high;
1153     long long pgsteal_normal;
1154     long long pgsteal_dma;
1155     long long pgscan_kswapd_high;
1156     long long pgscan_kswapd_normal;
1157     long long pgscan_kswapd_dma;
1158     long long pgscan_direct_high;
1159     long long pgscan_direct_normal;
1160     long long pgscan_direct_dma;
1161     long long pginodesteal;
1162     long long slabs_scanned;
1163     long long kswapd_steal;
1164     long long kswapd_inodesteal;
1165     long long pageoutrun;
1166     long long allocstall;
1167     long long pgrotated;
1168 };
```


Собираем статистику о виртуальной памяти

```
1589 long long get_vm_value(char *s)
1590 {
1591     int currline;
1592     int currchar;
1593     long long result = -1;
1594     char *check;
1595     int len;
1596     int found;
1597
1598     for (currline = 0; currline < proc[P_VMSTAT].lines; currline++) {
1599         len = strlen(s);
1600         for (currchar = 0, found = 1; currchar < len; currchar++) {
1601             if (proc[P_VMSTAT].line[currline][currchar] == 0 ||
1602                 s[currchar] != proc[P_VMSTAT].line[currline][currchar]) {
1603                 found = 0;
1604                 break;
1605             }
1606         }
1607         if (found && proc[P_VMSTAT].line[currline][currchar] == ' ') {
1608             result =
1609                 strtoll(&proc[P_VMSTAT].line[currline][currchar + 1],
1610                     &check, 10);
1611             if (*check == proc[P_VMSTAT].line[currline][currchar + 1]) {
1612                 fprintf(stderr, "%s has an unexpected format: >%s<\n",
1613                     proc[P_VMSTAT].filename,
1614                     proc[P_VMSTAT].line[currline]);
1615                 return -1;
1616             }
1617             return result;
1618         }
1619     }
1620     return -1;
1621 }
```

Функция read_vmstat() - записывает в vm_stat статистику о виртуальной памяти

```
1625 int read_vmstat()
1626 {
1627     proc_read(P_VMSTAT);
1628     if (proc[P_VMSTAT].read_this_interval == 0
1629         || proc[P_VMSTAT].lines == 0)
1630         return (-1);
1631
1632     /* Примечание:
1633     если запрошенная переменная не найдена
1634     в /proc/vmstat, тогда она установлена в -1 */
1635     GETVM(nr_dirty);
1636     GETVM(nr_writeback);
1637     GETVM(nr_unstable);
1638     GETVM(nr_page_table_pages);
1639     GETVM(nr_mapped);
1640     GETVM(nr_slab);
1641     GETVM(pgpgin);
1642     GETVM(pgpgout);
1643     GETVM(pswpin);
1644     GETVM(pswpout);
1645     GETVM(pgallocc_high);
```

```
1623 #define GETVM(variable) p->vm.variable = get_vm_value(__STRING(variable) );
```

int main(int argc, char **argv)

- занимает более 4000 строк кода
- “можно разделить на три секции”:
 - получение данных
 - вывод данных
 - обновление данных
- меньше всего содержит комментариев

```
"-----");  
"#   #   #   #####   #   #");  
"##   #   ##   ##   #   #   ##   #");  
"# #   #   #   ##   #   #   #   #   #");  
"#   #   #   #   #   #   #   #   #   #");  
"#   ##   #   #   #   #   #   ##");  
"#   #   #   #   #####   #   #");  
"-----");
```